ELEC3441: Computer Architecture
First Semester, 2026–27

Homework 2 (2026.2)

Due: See Moodle for the deadline

Instruction: Submit your answers electronically through Moodle.

There are 2 parts in this homework. Part A includes questions that aim to help you with understanding the lecture materials. They resemble the kind of questions you will encounter in quizzes and the final exam. Part A is an individual portion of this homework. You should be completing and turning in *your own work*.

Part B contains hands-on exercises that require you to design and evaluate processor systems using software tools, including the RISC-V compilation toolchain and ISA simulator. They are designed to help you understand memory-hierarchy performance and the use of evaluation tools. Part B is group work. Each group should turn in only 1 set of project files and 1 project report.

The following summarizes the 2 parts.

Part	Type	Indv/Grp
A	Basic problem set	Individual
B	Hands-on	Group of up to 2 persons

In all cases, you are encouraged to discuss the homework problems with classmates. However, you should not ask for or give out solution directly as that defeat the idea of having homework exercise. Giving out answers or copying answers directly will likely constitute an act of plagiarism.

Part A: Problem Set

A.1 Resolving Branch Hazards

You are designing the branch logic for a new deeply pipelined processor with 12 pipeline stages. Branches are resolved in stage 5 (stages count from 1). You are faced with two architectural choices to resolve control hazard:

s1. Speculate branch always-not-taken. Kill instructions if speculated wrong.

s2. Use a branch predictor. Your branch predictor is correct 95 % of the time.

Due to the increased hardware required for hazard detection and to kill the wrongly fetched instructions, the minimum cycle time of the processor is affected as follows:

S1	**S2**
1.5 ns	2 ns

Assume a single-issue pipeline with no structural hazards and sufficient forwarding to avoid data stalls in the shown sequence. Each misprediction costs 4 extra cycles. Under S1, every taken branch is a misprediction. Correct predictions cost none. Include 11 cycles for initial fill and final drain, and treat predictor accuracy as an expected value.

Now, consider the following code sequence for this question:

```
I1|       addi a1, zero, 10
I2|       addi a0, zero, 0
I3| loop: lw   a2, 0(a3)
I4|       lw   a4, 0(a5)
I5|       addi a3, a3, 4
I6|       addi a5, a5, 4
I7|       add  a2, a2, a4
I8|       srli a4, a2, 1
I9|       andi a4, a4, 1
I10|      beqz a4, skip
I11|      add  a0, a0, a2
I12|skip: nop
I13|      addi a1, a1, -1
I14|      bnez a1, loop
I15|      add  a7, zero, a0
I16|      addi a1, zero, 23
I17|      add  a0, a0, a1
```

A.1.1 Assume the array pointed to by `a3` and `a5` contains the following values:

```
int A[10] = { 2, 3, 5, 7, 11, 13, 17, 19, 23, 27}; // a3
int B[10] = { 3, 1, 4, 1,  5,  9,  2,  6,  5,  3}; // a5
```

Trace through the above code and record the sequence of branch and directions for the 2 branch instructions.

Sequence	**Instruction Address**	**Instruction**	**Branch Taken? (Y/N)**
1	I10	`beqz`	Y
2	I14	`bnez`	Y
3			
4	⋮	⋮	⋮

A.1.2 Consider solution S1 above and assume pipeline has full forwarding network and consider RAW data hazards only. How many cycles are needed to complete the above code segment? How many times does the processor need to kill wrongly fetched instructions?

A.1.3 Consider solution S2 above and assume pipeline has full forwarding network and consider RAW data hazards only. Assume your branch predictor has the same prediction accuracy for both branch instructions. How many cycles are needed to complete the above code segment?

A.1.4 Taking cycle time into consideration, which solution results in the best performance?

A.2 Cache Access

Consider the following sequence of memory accesses to the main memory in a 32-bit processor:

Address (hex)	Type
00000098	R
000001A8	W
00000288	R
000003B8	R
00000480	R
00000284	R
000003A0	R
000001A4	W
00000290	W
00000088	R

A.2.1 Assume the following data cache organization:

- Capacity: 512 B
- Line size: 8 words
- Organization: direct map
- Policy: write back, write allocate

Below is the current status of your cache:

Cache contents:

index	tag	valid	dirty	data							
				7	6	5	4	3	2	1	0
0	-	0	0	-	-	-	-	-	-	-	-
1	-	0	0	-	-	-	-	-	-	-	-
2	-	0	0	-	-	-	-	-	-	-	-
3	-	0	0	-	-	-	-	-	-	-	-
4	0	1	1	BEEF0007	BEEF0006	0EEE3441	0EEE3441	BEEF0003	BEEF0002	BEEF0001	0EEE3441
5	-	0	0	-	-	-	-	-	-	-	-
6	-	0	0	-	-	-	-	-	-	-	-
7	-	0	0	-	-	-	-	-	-	-	-
8	-	0	0	-	-	-	-	-	-	-	-
9	-	0	0	-	-	-	-	-	-	-	-
10	-	0	0	-	-	-	-	-	-	-	-
11	-	0	0	-	-	-	-	-	-	-	-
12	-	0	0	-	-	-	-	-	-	-	-
13	-	0	0	-	-	-	-	-	-	-	-
14	-	0	0	-	-	-	-	-	-	-	-
15	-	0	0	-	-	-	-	-	-	-	-

Each word is 4 bytes. The three writes in the trace store `0x11111111`, `0x22222222`, and `0x33333333`, respectively. Initially, all main-memory words contain `0x0EEE3441`; the dirty cache line may differ. Classify misses relative to an equally sized fully associative LRU cache with the same initial resident block. For the associative variants, retain the same resident memory block at `0x00000080`, with the displayed data, and rederive its tag and set index. All other lines start invalid. For the write-through variant, first write this initial dirty block back to memory and mark it clean, without counting this initialization as a trace access. Use LRU for the fully associative variant. Trace through the above memory access and answer the following:

(i) For each access, is it a hit or a miss?
(ii) Reason for miss (e.g. Conflict miss due to access X; Compulsory miss, etc)
(iii) Show the final content in the cache, including the tag.

A.2.2 Repeat A.2.1 but with a different cache organization:

- Organization: 2-way set associative
- Policy: true LRU, write back, write allocate

A.2.3 Repeat A.2.1 but with a different cache:

- Organization: Fully Associative
- Policy: write through, no write allocate

A.3 Simple Pipeline

In this question, you will examine the operation of pipelined processors using the following assembly code: The registers `a0` and `a1` initially point to input arrays A and B; `a2` and `a3` point to output arrays C and D. The arrays do not overlap and each contains at least 100 words. The stores write A[i]+B[i] to C[i] and A[i]-B[i] to D[i].

```
    # a0 <- 0x00A00000
    # a1 <- 0x00A10000
    add   s0, zero, a0
    add   s1, zero, a1
    add   s2, zero, a2
    add   s3, zero, a3
    addi  t0, zero, 100  # <-- N=100
L1: lw    t1, 0(s0)
    lw    t2, 0(s1)
    add   t3, t1, t2
    add   a0, a0, t3
    sw    t3, 0(s2)
    sub   t3, t1, t2
    add   a1, a1, t3
    sw    t3, 0(s3)
    addi  s0, s0, 4
    addi  s1, s1, 4
    addi  s2, s2, 4
    addi  s3, s3, 4
    addi  t0, t0, -1
    bne   t0, zero, L1
    add   a0, a0, a0
    add   a0, a0, a0
    add   a0, a0, a0
    add   a0, a0, a0
    add   a1, a1, a1
    add   a1, a1, a1
    add   a1, a1, a1
    add   a1, a1, a1
```

Consider a processor, P, which has the following properties:

- It has a pipeline with the following 7 stages:

```
IF1 IF2 ID EX MA1 MA2 WB
```

- It only has *one forwarding path from EX to ID*, i.e., only from output of the ALU at the end of EX stage to the pipeline register between ID and EX. Any data hazards that cannot be resolved by the forwarding path will cause the processor to stall while bubbles are inserted.
- A register written in WB can be read in ID in the same cycle. The EX-to-ID path can forward only the value produced by the instruction currently in EX; there are no forwarding paths from later pipeline registers. Store data and branch operands are read in ID. Instruction and data caches are separate, all accesses hit, and there are no structural hazards. Include pipeline fill and drain in the cycle count.
- Cache hit takes 2 cycles. Cache access is fully pipelined, so one new access can be requested every cycle.
- Branches are resolved in EX stage in parallel to the ALU. The processor speculates "always not-taken" on branches, i.e., the processor continues fetching and decoding instructions and kill the speculated instructions if branch is taken.

A.3.1 Assuming all elements in the array pointed to by `a0` and `a1` contain the *values* 100 and 17 respectively. Trace through the code. What is the value of `a0` and `a1` when the code fragment completes?

A.3.2 How many instructions are being *executed to completion* during run time?

A.3.3 List out all data hazards in the above code that would cause a stall in the pipeline.

A.3.4 Counting the number of instruction committing at WB stage, what is the average CPI of the above code?

A.4 Streaming Cache Performance

You are investigating the cache performance of your processor regarding the following code segment:

```
// int i, N, a, b;
// int y[], x[];
for (i = 0; i < N; i++) {
  y[i] = a*x[i] + b*y[i];
}
```

Now the above C code is compiled into the following RISC-V instructions:

```
# a0 initially contains the constant N
# a1 is base address of x[]
# a2 is base address of y[]
# constant a is stored in register a3
# constant b is stored in register a4
00| loop: lw    t1, 0(a1)
04|       lw    t2, 0(a2)
08|       mul  t0, t1, a3
0C|       mul  t1, t2, a4
10|       add   t0, t0, t1
14|       sw    t0, 0(a2)
18|       addi  a1, a1, 4
1C|       addi  a2, a2, 4
20|       addi  a0, a0, -1
24|       bne   a0, zero, loop
```

A.4.1 Cache Hit or Miss Assume array `x` is stored at memory address starting at `0xEA000000` while the array `y` is stored at memory address `0xEB000000`.

The processor data cache is initially empty and has the following properties:

- Capacity: 1 MiB
- Organization: 2-way set associative
- Line size: 8 words
- Replacement policy: true LRU
- Write policy: write through, no write allocate

A large write buffer is available such that the processor can resume running immediately after writing the data to the buffer.

Let $N = 2^{10}$, trace through the above code, then show and explain the sequence of cache hit or miss that will occur. Use the notation **RH** for read hit, **WH** for write hit, **RM** for read miss, and **WM** for write miss.

A.4.2 Based on your result above, what is the overall data cache miss rate of the above code when executed in the main CPU? Consider BOTH read and write access. Assume the write and read miss penalties are both 300 cycles. Hit time is 1 cycle. What is the average memory access time (AMAT) for this code?

A.4.3 For simplicity, assume there is no pipeline in the processor. Furthermore, assume `add`, `addi` and `bne` takes 1 cycle; `mul` takes 4 cycles; and performance of `lw` and `sw` depends on the cache performance. What is the total run time of the above code in cycles? You may leave the variable N in your answer.

A.4.4 Write Back Cache For the original $N = 2^{10}$ case, if the write policy of the cache is changed to *write back, with write allocate*, how would the following be affected? Consider 2 cases: (i) with write buffer, (ii) without write buffer.

(i) Read miss penalty
(ii) Write miss penalty
(iii) Overall performance of the above code

Explain your answer by tracing through the execution of the above code, highlighting any different in the required cache content handling with the use of write back cache.

A.5 Page Table Size

You are considering the design of the virtual memory system in your company's new 64-bit high performance embedded processor with the following features:

- 64-bit architecture
- 33-bit physical address

Use densely packed entries with the stated PTE size, and ignore alignment/padding overhead. Each first-level entry has the same size as a leaf PTE. Each second-level table is a fixed-size array: allocate the complete table when any entry in it is needed; leave wholly unused tables unallocated. Unless a part asks you to vary it, use $p = 12$. Calculate nominal page-table sizes even if they cannot fit in physical memory, and state when a proposed organization is infeasible.

Your system is hardwired with 8 GiB high performance memory. Each running process has its own page table to map their 64-bit virtual addresses to the physical memory locations. There is NO swap space on board. A process is killed with an out of memory error if it requests a page of memory from the OS when there is no physical space available.

The *maximum* memory address space allocated to each process is as follows. Depending on the need of the program, the process may not use up all the spaces that are allocated. Only used pages are allocated by the OS.

Segment	Address
code	`0x0000_0000_1000_0000 - 0x0000_0000_101F_FFFF`
stack	`0x0000_A000_0000_0000 - 0x0000_A000_0000_FFFF`
heap	`0xD000_0000_0000_0000 - 0xD000_0000_1FFF_FFFF`

A.5.1 Based on the above information, what is the maximum amount of memory (in bytes), that a process may use?

A.5.2 Let the size of a page be 2^p. What is the maximum number of pages that a process may be allocated in each segment?

A.5.3 In this part, consider the design of a simple 1-level *linear page table* that covers the entire 64-bit address space.

Let sizeof(PPN) denote the physical page number width in bits. The PTE size in bytes is:

$$\left\lceil \frac{\text{sizeof(PPN)} + 3}{8} \right\rceil$$

Assuming $p = 12$, what is the size of this linear page table for each process?

A.5.4 In order to reduce the total amount of storage required, you are considering to use a 2-level page table organization instead. In this scheme, the first level of page table is preallocated when a process is launched, while the second level is allocated on demand.

The first level of your page table contains 8192 entries. Given your page size is 2^p, how many PTEs are there in *each* second level page table?

A.5.5 Using your above 2-level page table, what is the maximum storage needed to implement all the necessary page tables if a process allocated all the possible pages?

A.5.6 Assume instead that a process only uses the *lower* 25 % of its address space in all 3 segments. In this case, what is the amount of memory needed to start the 2-level page table?

A.5.7 Given the above 25 % memory usage pattern, design an optimal 2-level page table system that would minimize the total size of page table. Assume $p = 12$.

A.5.8 As an OS designer, you have the choice of picking the location for the various areas. Suggest new scheme for locating the three segment, `code`, `stack` and `heap`, *together* with a design of 2-level page table system that may minimize the amount of storage needed for page table entries. If the location of the segments have no effect, please also explain why.

A.6 TLB Trace

In this exercise you will experiment with the interaction between the TLB and the page table in a VM system.

Assume the following system configuration:

- 16-bit virtual and 24-bit physical address
- 4 KiB page size
- 4-entry, *fully associative* TLB, true LRU.

Initially, the TLB and page table contains the following entries. An invalid entry is marked as "empty". An entry marked with disk$n$ has a page stored in hard disk, where n is the disk page number. Assume in this question that physical page number = disk page number + 0xD10.

TLB

Idx	VPN	PPN	Time
3	0x8	0x992	0028
2	0x2	0xB07	0011
1	0xA	0x289	0009
0	0xF	0x28A	0022

Page Table

VPN	PPN or Disk
0xF	0x28A
0xE	empty
0xD	empty
0xC	disk2
0xB	empty
0xA	0x289
0x9	empty
0x8	0x992
0x7	empty
0x6	empty
0x5	0x22E
0x4	empty
0x3	0xA9C
0x2	0xB07
0x1	empty
0x0	empty

A.6.1 Given the above VM setup, what is its TLB reach?

A.6.2 The following sequence of memory accesses are issued:

0x2468, 0x246C, 0x3070, 0xA500, 0x2470, 0x5CCC, 0xC000, 0x8004, 0x5CD0,

Answer the following:

1. For each memory access, is it a hit/miss in TLB, a hit/miss in the page table, a page fault?
2. What is the final state of the TLB and Page Table after the above accesses?

A.6.3 What are some of the advantages and disadvantages of using a larger page size?

A.6.4 You have the option to improve system performance by using an 8-entries direct map TLB. Assume the clock speed is the same for both cases, discuss what are the tradeoffs for using the new TLB compared to the original design. Give concrete examples on when (and if) one proposal will perform better than the other. State your assumptions.

A.7 TLB and Cache

based on a previous final exam question You are designing a paged memory system for a high-performance 32-bit processor for datacenter. The initial virtual memory subsystem design has the following characteristics:

- 32-bit virtual address
- 35-bit physical address
- 4 KiB page size
- Fully associative TLB with E entries, true LRU
- A single linear page table for each user process
- Page fault penalty: 10 000 ns
- On TLB hit, TLB access takes 0.5 ns
- Page tables are not cached; a page-table memory read takes 400 ns. A TLB lookup (hit or miss) takes 0.5 ns. Ignore TLB insertion time. Page fault service adds the stated penalty after the page-table lookup; it does not require another lookup in this model.
- Single level cache: *physically tagged*, direct map, 4 word line size, 1 MiB capacity,
- Cache hit time: 0.5 ns
- Cache miss penalty: 400 ns
- Clock rate: 1 GHz

A.7.1 Answer the following:

1	Number of bits needed for a virtual page number		bits
2	Number of bits needed for a physical page number		bits
3	Number of bits needed for page offset		bits

A.7.2 What is the TLB reach of the system?

A.7.3 If a TLB miss, the processor must reload the entry from the page table in memory. Given the above characteristic of this processor, how long does it take to update an entry to the TLB?

A.7.4 Assume your data cache always hit. After profiling a program B, you found that out of all the memory accesses, 3 % results in TLB **MISS**. Among the misses, 1 % results in a page fault. As a result of these address translations, what is the average memory access time (in seconds) for B?

A.7.5 You found that in general the data cache has a miss rate of 5 %. Your processor always complete address translations before accessing the data cache. That is, the system handles any TLB misses or page faults before accessing the cache. If a cache misses, the penalty is added on top of the time needed to translate an address.

Given this access mechanism, what is the average memory access time (AMAT) of the system, which includes both time needed for address translation and cache access?

Part B: Open-ended Project

B.1 Adapting Cache Characteristics

In this exercise, you will investigate and try to improve the cache performance of a set of benchmarks. To evaluate the benchmark performance, you will be using a RISC-V ISA simulator called `spike`. The `spike` simulator simulates the behavior of the RISC-V ISA with limited hardware implementation details. It features a built-in cache simulator that captures every memory access generated from the processor and collect statistics accordingly.

B.1.1 Getting Started The pinned RV32 C/C++ toolchain, `spike`, and the proxy kernel `pk` are included in the course Docker environment. Start the course container as described in the package README. Receive the homework release by following the update instructions in the course repository README.

In the container terminal, set the homework root and enter the benchmark directory:

```
export HW2ROOT=/workspace/homeworks/hw2
cd ${HW2ROOT}/benchmarks
```

The benchmark programs are each located in their own subdirectory.

B.1.2 Compiling Benchmark Programs The ISA simulator may execute any valid RISC-V program. If you examine the source code in each subdirectory, you will notice they are no different from any other normal program. Feel free to write your own benchmark if you are curious.

Compile all provided benchmarks using the top-level makefile:

```
cd ${HW2ROOT}/benchmarks
make
```

B.1.3 Running Simulator You may now execute your compiled program using the `spike` ISA simulator. Execute the target binary with an L1 instruction cache as follows:

```
cd ${HW2ROOT}/benchmarks/kmean
spike --ic=128:2:64 pk kmean
bbl loader
Verification passed!
I$ Bytes Read:            2501806104
I$ Bytes Written:         0
I$ Read Accesses:         625451526
I$ Write Accesses:        0
I$ Read Misses:           507
I$ Write Misses:          0
I$ Writebacks:            0
I$ Miss Rate:             0.000%
student$
```

The example counts above were recorded on 7 November 2026 using the pinned course environment with output displayed directly in the terminal. Use direct terminal execution consistently for all comparisons. Redirecting output to a file or pipe can change the counts, even with the same binary and tools. Compiler, proxy kernel, and `spike` versions can also affect counts. The displayed miss rate is rounded; use the raw access and miss counts in calculations.

In the above line, the last argument `kmean` specifies the RISC-V binary that you are simulating using `spike`. The argument `pk` before that stands for **p**roxy **k**ernel; it provides the small execution environment and proxies the program's system calls to the host.

Finally, the argument `-ic=128:2:64` tells `spike` to simulate an **i**nstruction **c**ache that has

- 128 sets
- 2-way set associative
- 64-byte block

If you multiply the three parameters, you get the capacity of the instruction cache as $128 \times 2 \times 64 = 16\,\text{KiB}$.

You can also specify the use of a data cache with the argument `-dc=<S>:<W>:<B>` and a unified L2 cache with the argument `-l2=<S>:<W>:<B>`. For example:

```
spike --ic=128:2:64 pk ...
spike --ic=128:2:64 --dc=128:2:64 pk ...
spike --ic=128:2:64 --dc=128:2:64 --l2=1024:4:64 pk ...
```

B.1.4 Cache Evaluation Now, run all the benchmark programs with the following 2 memory hierarchies, one with L2 and the other without L2 cache:

- `-ic=128:1:32 -dc=256:1:32`
- `-ic=128:1:32 -dc=256:1:32 -l2=1024:2:128`

Collect statistics about the memory hierarchies from the output and answer the following questions.

1. For each benchmark, what is the miss rate for L1 I$, L1 D$ and L2$? Which benchmark has the best and which one has the worst cache performance? *Hint:* Consider automating the process, as you will probably need to regenerate a lot of similar statistics in the rest of this homework.
2. What is the cache access time for the two L1 caches according to Table B.1?
3. Based on your above answer, what is the cycle time of the pipeline? Assuming on a cache hit, data should be returned in 1 cycle. Also assume that critical path of all non-memory pipeline stages

is 600 ps. In other word, your cycle time is limited by the cache if the cache access time is longer than 600 ps. Otherwise, cache access is not the bottleneck.

4. Calculate the average CPI for the benchmarks *without L2 cache*. You can use the following formula to calculate average CPI, where MP = Miss Penalty, CT = Cycle Time. Use $\text{CPI}_{\text{base}} = 1.2$.

   Assume the backside of L1 caches are connected to a single DRAM with 100 ns access time, i.e., accessing the first word from the DRAM takes 100 ns. Each subsequent word takes additional 5 ns.

   $$\text{CPI} = \text{CPI}_{\text{base}} + \frac{\#\text{L1_I\$ misses} + \#\text{L1_D\$ misses}}{\#\text{instructions}} \times \left\lceil \frac{\text{MP}}{\text{CT}} \right\rceil$$

   Note that the number of instruction is the `I$ Read Accesses` number in the output of `spike`.

5. What is the AMAT (in ns) of the L2$ for all benchmarks? Use the following formula to calculate AMAT. (HT=Hit Time, MR=Miss Rate, MP=Miss Penalty) Assume the backside of the L2$ is connected to a DRAM with 100 ns access time. Note that L2 AMAT is the time calculated after L1 access. So HT is the time to access L2 cache after L1, MR is the local L2 miss rate, etc. Assume L2 access time is 10 ns. In this exercise a word is 4 bytes. Fill the entire cache line before resuming: use $100 + (w - 1) \times 5$ ns for a line of w words, including the 128-byte L2 line. Ignore dirty-eviction latency and overlap; this is an analytical model, not a cycle-accurate Spike result.

   $$\text{AMAT}_{L2} = \text{HT}_{L2} + \text{MR}_{L2} \times \text{MP}_{L2}$$

6. Calculate the average CPI for the benchmarks with L2 cache. Use the following formula to calculate CPI, assuming the L2$ is running asynchronously on its own clock domain. Use the same base CPI as above.

   $$\text{CPI} = \text{CPI}_{\text{base}} + \frac{\#\text{L1_I\$ misses} + \#\text{L1_D\$ misses}}{\#\text{instructions}} \times \left\lceil \frac{\text{AMAT}_{L2}}{\text{CT}} \right\rceil$$

7. Based on your answers above, does the L2$ help with performance?

B.1.5 Design Space Exploration for L1 Data Cache Configuration Use the no-L2 baseline with `-dc=256:1:32`, with a perfect I-cache for both baseline and candidates.

Restrict capacities to the columns in Table B.1 and exclude N/A cells. A word is 4 bytes; pass line sizes to Spike in bytes. Keep benchmark inputs fixed. Compute modeled runtime as instruction count times modeled CPI times cycle time, and maximize the geometric mean of baseline runtime divided by candidate runtime.

Find an optimal L1 D$ configuration that maximizes performance of the benchmarks. You can assume L1 I$ is perfect and does not affect processor performance. There is no L2$. Pick one configuration from the following design space:

- Total capacity: up to 2 MiB
- Associativity: $1, 2, 4, 8, 16, 32, 64$
- Cache line size: 4, 8 (words)

Note the following:

- Your cache organization may affect your processor cycle time. Your architecture requires that L1 cache returns data in 1 cycle.
- Use geometric mean of normalized performance of the benchmark programs as the overall performance metric.

Hint: You should write a program or use a spreadsheet program like Excel to help you find the optimal L1 configuration. Remember, you need to compute the geometric mean of speed up of different configurations relative to the baseline.

B.1.6 Optional: Improved Benchmarks Now, given your optimal L1 cache obtained from above, try to implement a faster version of the *slowest* benchmark program than the provided implementation. You may use any modifications to data structure or code, as long as the same calculation is still performed at runtime, i.e., you cannot simply hard code the answer at compile time. Evaluate the improvements you make using the geometric mean of speedups that you achieve over the baseline that you determined in the previous question.

B.1.7 Submission Submit your answer to B.1.4, B.1.5, and the optional part in B.1.6. Make your answer concise, but make sure you have data to support your analysis.

(a) cache line size = 4 words

assoc \ size	4KB	8KB	16KB	32KB	64KB	128KB	256KB	512KB	1MB	2MB
1	0.21	0.27	0.32	0.38	0.46	0.54	0.69	0.85	1.14	1.35
2	0.37	0.44	0.50	0.54	0.57	0.64	0.73	1.00	1.24	1.51
4	0.44	0.49	0.52	0.57	0.61	0.69	0.79	1.02	1.31	1.58
8	N/A	0.67	0.70	0.75	0.77	0.83	0.90	1.18	1.61	1.98
16	N/A	N/A	1.05	1.08	1.13	1.17	1.24	1.46	2.43	2.74
32	N/A	N/A	N/A	1.83	1.84	1.91	2.00	2.44	2.85	4.97
64	N/A	N/A	N/A	N/A	3.39	3.67	3.87	4.25	4.72	6.06

(b) cache line size = 8 words

assoc \ size	4KB	8KB	16KB	32KB	64KB	128KB	256KB	512KB	1MB	2MB
1	0.24	0.28	0.30	0.36	0.47	0.54	0.70	0.92	1.15	1.62
2	0.51	0.53	0.54	0.56	0.60	0.65	0.75	0.96	1.27	1.80
4	N/A	0.73	0.74	0.76	0.80	0.83	0.87	1.07	1.44	2.21
8	N/A	N/A	1.13	1.14	1.18	1.21	1.24	1.66	2.24	3.14
16	N/A	N/A	N/A	1.80	1.94	1.97	1.98	2.63	3.34	4.47
32	N/A	N/A	N/A	N/A	2.55	2.37	3.56	3.85	4.65	6.35
64	N/A	N/A	N/A	N/A	N/A	4.50	5.23	5.78	6.47	7.20

Table B.1: Cache access time (in ns) for various cache configurations.